

PREPARED FOR GOODYEAR

What Makes Glean Different

One secure, permission-aware intelligence layer across teams — from Goodyear Proprietary Systems to enterprise-wide transformation.

Glean · Enterprise AI Platform

THE TRANSFORMATION OUTCOME

The goal isn't another AI tool. It's a new way to work.

Goodyear is moving away from long, siloed cycles toward agile, discovery-led ways of working. The outcome we're after is one secure intelligence layer across teams — not another point tool.

Today — fragmented

Separate tools per team (Copilot, Claude, ChatGPT, internal builds); knowledge trapped in silos, so answers see only a slice of the data.

The outcome — one intelligence layer

One secure, permission-aware layer across teams — from isolated use cases to enterprise transformation.

WHERE IT STARTS · GOODYEAR PROPRIETARY SYSTEMS

One problem that becomes the blueprint

Goodyear Proprietary Systems helps tire-development teams work from trusted, peer-reviewed, proprietary knowledge. It's the entry point that scales:

1 → ∞

One source today → add sources with no rebuild

+1 dept

Extend the same layer to an adjacent team

Cross-fn

Answers span product, supply chain, creative, finance

Scale

The blueprint repeats wherever context matters

WHY GLEAN IS DIFFERENT

Five differences that turn AI into transformation

Starts with the outcome

Glean begins with the business problem and how you'll measure change — cost, speed, quality, risk — not with a tool.

Horizontal across the enterprise

One context layer over people, content, and systems — so a product question can reason across engineering, supply chain, and finance.

Complements your models

Model-agnostic. Glean makes Copilot, Claude, and others more useful — no rip and replace.

Secure & permission-aware

Inherits existing access controls, so users see only what they're allowed to. Proprietary tire knowledge stays protected.

Cheap to scale from 1 to many

Prebuilt connectors, governance, and reusable agents mean expanding to new teams doesn't mean rebuilding.

Context is the moat

Answers grounded in your enterprise are preferred 2.5x over generic model output — and cut token spend by reaching data first.

PROVEN WHERE THE WORK IS PHYSICAL

Automotive & manufacturing peers already scaled this

140K

GM employees live on Glean — company-wide in ~3 months

\$900M

Identified savings at GM, with 15.8x 3-yr ROI


25K

Nationwide users — pilot to full rollout in 8 weeks

90%+


Typical adoption once teams start using Glean

Adoption scales fast, not for years



GM rollout — users on Glean (pilot → month 3)

Consumption cost comes down



Reaching enterprise data before an LLM + query optimization (% of baseline)

THE PATH FORWARD

Measurable outcomes in weeks, not years

1 · Blueprint

Align on outcome, priority use cases, success criteria

2 · Demo

Walk the Goodyear team through your real workflows

3 · Pilot

Expand Goodyear Proprietary Systems beyond one source; measure change

4 · Scale

Reuse the same governed layer across lines of business

The real question isn't whether Glean can support Goodyear Proprietary Systems — it's whether it becomes the first proof of a broader, governed way for product, engineering, and enterprise teams to work.

Turn scattered research into validated product decisions

The Product COE is pushing Goodyear from long waterfall cycles toward agile, discovery-led work. Glean keeps every step grounded in Goodyear's own research, calls, and prior decisions — so discovery is faster and better evidenced.

Interactive PM workflow · From customer signal to validated PRD

End-to-end time today: **390 min**

Glean cuts this to 75 min · 81% faster

Every step stays grounded in Goodyear's own research, calls, and prior decisions.

Input	Value	
PMs & designers on the team	25	\$1.07M Value of time returned / year
Workflows per person / week	3	
Hours per workflow today	6	11,880 Hours returned / year
Time reduced with Glean	55%	
Fully-loaded hourly rate	\$90	6.2 FTE-equivalent capacity unlocked
		\$3.21M Value over 3 years

Give engineers the full picture in one place

Engineers lose time stitching context across tools to diagnose and resolve issues. Glean pulls errors, telemetry, prior fixes, and code into one place, then drafts the PR inside your existing GitHub workflow.

Interactive PM workflow · Diagnose → resolve → document an issue

End-to-end time today: **285 min**

Glean cuts this to 60 min · 79% faster

Glean pulls errors, telemetry, prior fixes, and code into one place, then drafts the PR in your existing workflow.

Input	Value	
Engineers on the team	120	\$4.41M Value of time returned / year
Workflows per person / week	5	
Hours per workflow today	4	51,840 Hours returned / year
Time reduced with Glean	45%	
Fully-loaded hourly rate	\$85	27.0 FTE-equivalent capacity unlocked
		\$13.22M Value over 3 years

One governed layer instead of AI sprawl

Goodyear's AI strategy is fragmented across thousands of applications, with rising consumption and no standard governance. Turnkey, permission-aware connectors inherit your existing access — most of the setup effort disappears.

Interactive PM workflow · Stand up a governed AI use case

End-to-end time today: **24 wks**

Glean cuts this to 5.0 wks · 79% faster

Turnkey, permission-aware connectors inherit your existing access — most of the effort disappears.

Input	Value
Employees using AI	2,200
Avg AI spend / user / month	\$45
Consumption reduction with Glean	40%
Redundant point tools retired	4
Avg annual cost per retired tool	\$200K

\$1.28M

Total savings / year

\$475K

AI consumption saved / year

\$800K

Tool consolidation savings / year

\$3.83M

Total savings over 3 years

What it is, and how it works

The “Tao” is simply the way Glean approaches enterprise AI. Instead of bolting on tools, you engineer enterprise context once and reuse it everywhere.

Engineer context once

Connect data, permissions, the knowledge graph, and memory a single time — then reuse it everywhere.


Stay open & secure

Model-agnostic and permission-aware end to end, so you use any model without lock-in and deploy without fear.

Grow with maturity


Start with Search, add the Assistant, package Skills, then run Agents — the same foundation carries each step.

Why the Glean way wins



Glean vs. a Microsoft-centric point tool vs. a DIY internal build

Capability grows on one foundation



Autonomy increases; the context layer stays the same

Context

Answers preferred 2.5x over generic output

Open

Multi-model — no lock-in

Horizontal

Build once, deploy everywhere

Secure

Permission-aware by default

Loved

90%+ adoption is typical

Confidence

Transformation you can trust


More outcome per token — and materially lower AI spend

Consumption climbs when context is missing: teams burn tokens on every query, then re-run calls to fix brittle answers. Glean cleans, de-duplicates, and structures messy natural language — and does the same inside agents — so the model does less work for a better answer.

The language rewrite in action

Example	Raw input	Glean-harmonized	Tokens
Airtable update	"hey can you please go update the goodyear innovation tracker in airtable with everything we talked about in today's meetings and make sure the statuses are right"	<code>update(tracker="Goodyear Innovation", source=today.meetings, fields=["status"])</code>	44→33 -25%
Research question	"can you find me all the stuff we have on the new compound testing and summarize what the results were so far please"	<code>search(topic="new compound testing") -> summarize(scope="results")</code>	40→30 -25%
Agent instruction	"every morning build me a dashboard that shows what i need to do, who to follow up with, and where adoption might be slipping"	<code>agent.daily: build(dashboard=["tasks", "followups", "adoption_risk"])</code>	42→31 -26%

Where the savings come from



Cumulative AI consumption cost vs. baseline

~30%

Saved by reaching enterprise data before the LLM

+25%

More from rewriting query/agent language into code

up to 50%

Total AI consumption savings seen across clients


2.5x

Better answers from context — fewer costly retries

Glean vs. Copilot Premium

Copilot Premium is strong inside the Microsoft ecosystem. But Goodyear's work spans far beyond Microsoft — engineering, supply chain, creative, and proprietary systems. Glean is the enterprise-wide intelligence layer that unifies all of it.

Fit across what Goodyear actually runs on



Relative coverage & fit by capability (illustrative)

How they compare, capability by capability

Capability	Glean	Copilot Premium
Source coverage	150+ connectors across Microsoft and non-Microsoft — ServiceNow, Workday, GitHub, Airtable, Salesforce, plus proprietary systems like tire-dev knowledge.	Best within the Microsoft Graph; non-Microsoft sources are limited or require extra work.
Model choice	Model-agnostic. Use GPT, Claude, Gemini — even Copilot — and switch as the market evolves. No lock-in.	Tied to Microsoft's chosen models.
Cross-app reasoning	A knowledge graph connects people, content, and activity so answers reason across departments and systems.	Strongest reasoning over Microsoft 365 content; cross-system context is thinner.
Custom agents	Build once, reuse everywhere — agents and skills that run across every connected source and team.	Copilot Studio agents are capable but Microsoft-centric.
Security & permissions	Permission-aware end to end, inheriting existing access from every source system.	Enterprise-grade within Microsoft's permission model.
Adoption	90%+ adoption is typical; one horizontal experience for the whole company.	Adoption often concentrated among heavy Microsoft 365 users.

Capability	Glean	Copilot Premium
Cost efficiency	Reaches enterprise data before an LLM and rewrites queries into code — up to 50% lower AI consumption cost.	Fixed per-seat premium; consumption not optimized across sources.
Relationship	Complements Copilot — Glean is the context and governance layer that makes every model, including Copilot, more useful.	One assistant, one ecosystem.

Bottom line: this isn't Glean or Copilot. Glean is the enterprise-wide intelligence layer that unifies Goodyear's full stack and makes whatever models you use — Copilot included — deliver better, cheaper answers.